

# ECE433/COS435 Introduction to RL

## Assignment 4: Q-learning, DQN, and Approximate Dynamic Programming

### Spring 2024

Fill me in

Your name here.

Due March 4, 2024

## Collaborators

Fill me in

Please fill in the names and NetIDs of your collaborators in this section.

## Instructions

Writeups should be typeset in Latex and submitted as PDFs. You can work with whatever tool you like for the code, but **please submit the asked-for snippet and answer in the solutions box as part of your writeup. We will only be grading your write-up.** Make sure still also to attach your notebook/code with your submission.

## Question 1. Q Learning

### Tabular setting

If the state and action spaces are sufficiently small, we can simply maintain a table containing the value of  $Q(s, a)$  – an estimate of  $Q^*(s, a)$ <sup>1</sup> – for every  $(s, a)$  pair. In this *tabular setting*, given an experience sample  $(s, a, r, s')$ , the update rule is

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left( r + \gamma \max_{a' \in \mathcal{A}} Q(s', a') - Q(s, a) \right) \quad (1)$$

where  $\alpha > 0$  is the learning rate and  $\gamma \in [0, 1]$  is the discount factor.

---

<sup>1</sup>Here,  $Q^*(s, a)$  refers to optimal  $Q$  value function.

## Question 1.a: Regular Q-Learning

Why is it difficult to extend this learning rule to the game of Tetris or similar Atari games?

### Solution

Due to the scale of Atari environments, we cannot reasonably learn and store a  $Q$  value for each state-action tuple.

## Approximation setting

Here, we instead represent our  $Q$  values as a function  $\hat{q}(s, a; \mathbf{w})$ , where  $\mathbf{w}$  are parameters of the function (typically a neural network's weights and bias parameters). In this *approximation setting*, the update rule becomes

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha \left( r + \gamma \max_{a' \in \mathcal{A}} \hat{q}(s', a'; \mathbf{w}) - \hat{q}(s, a; \mathbf{w}) \right) \nabla_{\mathbf{w}} \hat{q}(s, a; \mathbf{w}). \quad (2)$$

In other words, given current parameters at iteration  $i$ ,  $\mathbf{w}_i$ , then we aim to minimize the loss at  $\mathbf{w}_{i+1}$  which is:

$$L(\mathbf{w}_{i+1}) = \mathbb{E}_{s,a,r,s' \sim \mathcal{D}} \left[ \left( r + \gamma \max_{a' \in \mathcal{A}} \hat{q}(s', a'; \mathbf{w}_i) - \hat{q}(s, a; \mathbf{w}_{i+1}) \right)^2 \right] \quad (3)$$

## Question 1.b: Action spaces

We can represent a Q-function  $Q(s, a)$  as either a function  $Q(s; w) : S \rightarrow R^A$  outputting the vector of Q-values  $[Q(s, a_1), \dots, Q(s, a_{|A|})]$  all at once, or a function  $Q(s, a; w) : S \times A \rightarrow R$  outputting a single Q-value  $Q(s, a)$ . What is a benefit of implementing the former over the latter? What is one drawback?

### Solution

A benefit of the former representation over the latter is that it is easier to compute the max over actions with just one forward pass. A drawback of the former representation over the latter is that by representing the output as a power of  $|A|$ , the Q function is hard to scale in a setting with many actions.

## Question 1.c: Continuous actions

Consider an environment such as Mountain Car Continuous where the action space is  $[-1, 1] \in \mathbb{R}$ . How might our representation of the Q-function described in (1.a) change?

### Solution

In environments where the action space is discrete, the representation of the Q-function involves taking the maximum over Q-values corresponding to discrete actions. In a continuous action space, you cannot simply enumerate all possible actions and take the maximum Q-value because there are infinitely many possible actions. Hence, instead of maintaining a table of Q-values for each state-action pair, we need to use function approximation to represent the Q-function. A common approach is to use neural networks to approximate the Q-function. These networks can take a state and action as input and output a Q-value, thereby allowing us to work with continuous action spaces.

### Question 1.d: Policy iteration

Policy iteration is a model-based (i.e. we have access to the transition probabilities) reinforcement learning algorithm that provably improves a policy. In policy iteration, there is a step called "policy evaluation" that estimates the "value" of a state under the current policy  $\pi$  being learned:

$$V^\pi(s) = \sum_{s' \in S} P_{\pi(s)}(s' | s) [r(s, a, s') + \gamma V^\pi(s')]$$

The Q-learning update as described in Equation 1, on the other hand, models an entirely different value function. Other than the fact Q-learning learns a Q function and policy evaluation learns a V function, how do these learned value functions differ from each other and why?

### Solution

The main difference between the two methods is that policy iteration requires a model to compute the full expected value of a policy, while Q-learning updates its estimates based only on the observed rewards and the maximum value of the next state. Policy iteration systematically evaluates and improves a policy (learning  $V^\pi$ ), whereas Q-learning directly learns the value of the optimal policy ( $V^*$ ).

The reason for this difference is due to the max operator used in Q-learning. This max operator is what ensures that Q-learning is learning the optimal value function  $V^*$  as it always backs up the value of the best possible future action. This is opposed to policy evaluation in policy iteration, which only considers the expected value under the current policy  $\pi$  without necessarily looking for the maximum value.

### Question 1.e: Two Learning Rules for Q-Values

Assuming that you're given a dataset of transitions  $(s_t, a, r, s_{t+1})$  collected from a policy  $\beta(a|s)$ . Consider learning a Q-function using one of these two learning rules:

#### Learning Rule 1

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left( R_{t+1} + \gamma \max_{a \in \mathcal{A}} Q(s_{t+1}, a) - Q(s_t, a_t) \right) \quad (4)$$

## Learning Rule 2

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left( R_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t) \right) \quad (5)$$

Assume that states and actions are discrete, so the Q-function is just a table. In what settings will these two learning rules converge to the same policy?

### Solution

Q-learning will converge to the optimal policy as long as all state-action pairs are visited an infinite number of times and the learning rate  $\alpha$  decreases over time according to certain conditions. The optimal policy that Q-learning converges to is one that always selects the action with the highest Q-value in every state.

SARSA will converge to the optimal policy under the same conditions as Q-learning if the policy used to update the Q-values is greedy. However, if the policy used is not greedy then SARSA will converge to a near-optimal policy that is influenced by the level of exploration specified by  $\epsilon$ .

These converge to similar values when, as described above, SARSA follows a greedy policy and thus effectively becomes Q-learning, as the action taken will always be the one that maximizes the Q-value in the next state (or more specifically SARSA is used with a decay schedule for  $\epsilon$  that goes to zero, ensuring that the policy becomes greedy in the limit).